

# Forecasting at Scale

Christian Barz

2020-07-23

1 Prophet's goals and alternative

2 Prophet model

3 Stan model and fitting

4 Future Talk's

5 Further reading

## Prophet's goals and alternative

# Prophet library

The Prophet library<sup>1</sup> is a model and a framework, that targets

- non-experts with background business knowledge

---

<sup>1</sup>Taylor, Letham, Forecasting at scale, 2017, The American Statistician

# Prophet library

The Prophet library<sup>1</sup> is a model and a framework, that targets

- non-experts with background business knowledge
- less knowledge about time series is required

---

<sup>1</sup>Taylor, Letham, Forecasting at scale, 2017, The American Statistician

# Prophet library

The Prophet library<sup>1</sup> is a model and a framework, that targets

- non-experts with background business knowledge
- less knowledge about time series is required
- easy to configure (analyst in the loop)

---

<sup>1</sup>Taylor, Letham, Forecasting at scale, 2017, The American Statistician

# Prophet library

The Prophet library<sup>1</sup> is a model and a framework, that targets

- non-experts with background business knowledge
- less knowledge about time series is required
- easy to configure (analyst in the loop)
- easily interpretable parameters

---

<sup>1</sup>Taylor, Letham, Forecasting at scale, 2017, The American Statistician

# Prophet library

The Prophet library<sup>1</sup> is a model and a framework, that targets

- non-experts with background business knowledge
- less knowledge about time series is required
- easy to configure (analyst in the loop)
- easily interpretable parameters
- flexible for a wide range of business problems

---

<sup>1</sup>Taylor, Letham, Forecasting at scale, 2017, The American Statistician



# Alternative

The forecast R-package contains some alternative for automated times series models.

- `auto.arima`<sup>2</sup>, fits multiple ARIMA models and take the best fit

---

<sup>2</sup>Hyndman, Khandakar et al. 2007, Automatic time series for forecasting: the forecast package for R

# Alternative

The forecast R-package contains some alternative for automated times series models.

- `auto.arima`<sup>2</sup>, fits multiple ARIMA models and take the best fit
- `ets`<sup>3</sup>, fits multiple exponential smoothing models and take the best fit

---

<sup>2</sup>Hyndman, Khandakar et al. 2007, Automatic time series for forecasting: the forecast package for R

<sup>3</sup>Hyndman, Koehler., Snyder & Grose, 2002, 'A state space framework for automatic forecasting using exponential smoothing methods'

# Alternative

The forecast R-package contains some alternative for automated times series models.

- `auto.arima`<sup>2</sup>, fits mutliple ARIMA models and take the best fit
- `ets`<sup>3</sup>, fits mutliple exponential smoothing models and take the best fit
- `snaive`<sup>4</sup>, random walk model with seasonality

---

<sup>2</sup>Hyndman, Khandakar et al. 2007, Automatic time series for forecasting: the forecast package for R

<sup>3</sup>Hyndman, Koehler., Snyder & Grose, 2002, 'A state space framework for automatic forecasting using exponential smoothing methods'

<sup>4</sup>De Livera, Hyndman. & Snyder, 2011, 'A state space framework for automatic forecasting using exponential smoothing methods'

# Prophet model

# Generalized additive models

Prophet uses a generalized additive model (GAM)<sup>5</sup>

$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$

- $y(t)$  target

---

<sup>5</sup>Hastie & Tibshirani, 1987, 'Generalized additive models: some applications'

# Generalized additive models

Prophet uses a generalized additive model (GAM)<sup>5</sup>

$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$

- $y(t)$  target
- $g(t)$  trend function

---

<sup>5</sup>Hastie & Tibshirani, 1987, 'Generalized additive models: some applications'

# Generalized additive models

Prophet uses a generalized additive model (GAM)<sup>5</sup>

$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$

- $y(t)$  target
- $g(t)$  trend function
- $s(t)$  Fourier series for periodic changes

---

<sup>5</sup>Hastie & Tibshirani, 1987, 'Generalized additive models: some applications'

# Generalized additive models

Prophet uses a generalized additive model (GAM)<sup>5</sup>

$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$

- $y(t)$  target
- $g(t)$  trend function
- $s(t)$  Fourier series for periodic changes
- $h(t)$  represents holiday effects, i.e. irregular schedules over one or more days

---

<sup>5</sup>Hastie & Tibshirani, 1987, 'Generalized additive models: some applications'



# Generalized additive models

Prophet uses a generalized additive model (GAM)<sup>5</sup>

$$y(t) = g(t) + s(t) + h(t) + \epsilon_t$$

- $y(t)$  target
- $g(t)$  trend function
- $s(t)$  Fourier series for periodic changes
- $h(t)$  represents holiday effects, i.e. irregular schedules over one or more days
- $\epsilon_t$  error term (not accommodated by the model)

---

<sup>5</sup>Hastie & Tibshirani, 1987, 'Generalized additive models: some applications'

# The trend model

Trend can be modeled with

- saturated growth model

# The trend model

Trend can be modeled with

- saturated growth model
- piecewise linear model

# Nonlinear Saturating growth (basic)

$$g(t) = \frac{C}{1 + \exp(-k(t - m))}$$

- $C$  carrying capacity (upper bound)

# Nonlinear Saturating growth (basic)

$$g(t) = \frac{C}{1 + \exp(-k(t - m))}$$

- $C$  carrying capacity (upper bound)
- $k$  growth rate

# Nonlinear Saturating growth (basic)

$$g(t) = \frac{C}{1 + \exp(-k(t - m))}$$

- $C$  carrying capacity (upper bound)
- $k$  growth rate
- $m$  offset parameter

# Nonlinear Saturating growth (time dependent parameters)

- $C$  (carrying capacity) and  $k$  (growth rate) are usually not constant

# Nonlinear Saturating growth (time dependent parameters)

- $C$  (carrying capacity) and  $k$  (growth rate) are usually not constant
- incorporate trend changes by explicit defined change points



# Nonlinear Saturating growth (time dependent parameters)

- $C$  (carrying capacity) and  $k$  (growth rate) are usually not constant
- incorporate trend changes by explicit defined change points
- change point can be set by analyst (e.g. product launches) or automatically selected

# Nonlinear Saturating growth (time dependent carrying capacity)

TBA - this is not explained in the non reviewed paper

# Nonlinear Saturating growth (non constant growth rate - idea)

let  $s_1, \dots, s_n$  be change points (time stamps).

define  $\delta \in \mathbb{R}^n$  (vector of rate adjustments)

define growth rate at time  $t$  by

$$k(t) := k + \sum_{j:s_j < t} \delta_j$$

where  $k$  is the base growth rate

# Nonlinear Saturating growth (non constant growth rate - mathematically correct)

define  $a(t) \in \{0, 1\}^n$  by

$$a_j(t) := \begin{cases} 1, & \text{for } t \leq s_j \\ 0, & \text{otherwise} \end{cases}$$

then rate at time  $t$  is given by

$$k(t) := a(t)^T \delta$$

# Nonlinear Saturating growth (non constant growth rate - mathematically correct)

The offset parameter  $m$  is then adjusted according to the change points by a formula:

$$\gamma := f(s, m, k, \delta).$$

Putting all together we receive :

$$g(t) = \frac{C(t)}{1 + \exp(-(k + a(t)^T \delta)(t - (m + a(t)^T \gamma)))}$$

# Linear trend with Changepoints

using the change points from above a piecewise linear model is given by

$$g(t) := (k + a(t)^T \delta) * t + (m + a(t)^T \gamma)$$

# Automatic changepoint selection

Change points are automatically detected by putting a sparse prior on  $\delta$ , i.e.

$$\delta_j \sim \text{Laplace}(0, \tau)$$




Figure 1: LaPlace density function from wikipedia, here  $\sigma = \delta$

# The seasonal model

The seasonality is approximated by a (truncated) Fourier serie:

$$s(t) = \sum_{i=1}^N (a_n \cos(\frac{2\pi it}{P}) + b_n \frac{2\pi it}{P})$$

where

- P is the period in days, for example  $P = 7$  for weekly seasonality



# The seasonal model

The seasonality is approximated by a (truncated) Fourier serie:

$$s(t) = \sum_{i=1}^N (a_n \cos(\frac{2\pi it}{P}) + b_n \frac{2\pi it}{P})$$

where

- P is the period in days, for example  $P = 7$  for weekly seasonality
- N truncates the series and is a model parameter for the user to adjust the fitting

# The seasonal model (fitting)

Fitting requires to estimate  $2N$  parameters  $\beta = (a_1, b_1, \dots, a_N, b_N)$ .

This is done by constructing a matrix of seasonality vectors for each  $t$  in our historical and future data, e.g. for weekly seasonality and  $N = 3$

$$X(t) := (\cos(\frac{2\pi(1)t}{7}), \dots, \sin(\frac{2\pi(3)t}{7}))$$

Hence the seasonal component is

$$s(t)X(t) \cdot \beta$$

where the prior on  $\beta$  is normally distributed:  $\beta \sim \text{Normal}(0, \sigma^2)$ .

# Fourier series

TBF

- recall how is an element in an vector space is represented in a basis

# Fourier series

## TBF

- recall how is an element in an vector space is represented in a basis
- the functions  $e_k(x) := \exp(\frac{2\pi i k x}{P})$  is a basis of a (dense) subspace in the vector space of  $P$  periodic functions

# Fourier series

## TBF

- recall how is an element in an vector space is represented in a basis
- the functions  $e_k(x) := \exp(\frac{2\pi i k x}{P})$  is a basis of a (dense) subspace in the vector space of  $P$  periodic functions
- the basis functions  $e_k$  can be expressed as a function in sin and cos

# Fourier series

## TBF

- recall how is an element in an vector space is represented in a basis
- the functions  $e_k(x) := \exp(\frac{2\pi i k x}{P})$  is a basis of a (dense) subspace in the vector space of  $P$  periodic functions
- the basis functions  $e_k$  can be expressed as a function in sin and cos
- convergence of Fourier series

# “Holidays and Events”

TBF

## Stan model and fitting



# Stan model and fitting

Stan is a platform for Bayesian inference using MCMC sampling and more.

# Stan model with logistic growth trend model

```
model {  
  // Priors  
  k ~ normal(0, 5);  
  m ~ normal(0, 5);  
  epsilon ~ normal(0, 0.5);  
  delta ~ double_exponential(0, tau);  
  beta ~ normal(0, sigma);  
  // Logistic likelihood  
  y ~ normal(C ./ (1 + exp(-(k + A * delta)  
                          .* (t - (m + A * gamma)))) +  
            X * beta, epsilon);  
}
```

# Stan fitting for prophet

- Prophet uses the L-BFGS algorithm from Stan to fit the GAM

# Stan fitting for prophet

- Prophet uses the L-BFGS algorithm from Stan to fit the GAM
- BFGS = Broyden–Fletcher–Goldfarb–Shanno algorithm, quasi Newton method

## Future Talk's

# regarding Prophet

- how the time dependent carrying capacity is determined ?

# regarding Prophet

- how the time dependent carrying capacity is determined ?
- how additional regressors are included ?

# regarding Prophet

- how the time dependent carrying capacity is determined ?
- how additional regressors are included ?
- how outliers are handled ?



# theoretical background

## ■ Stan models

# theoretical background

- Stan models
- GAMS

# theoretical background

- Stan models
- GAMS
- Bayesian models

# theoretical background

- Stan models
- GAMS
- Bayesian models
- L-BFGS

# theoretical background

- Stan models
- GAMS
- Bayesian models
- L-BFGS
- Marcov chain Monte Carlo

## Further reading

## Further reading

- The textbook for forecast R-package, but with a lot of theory and practice